

Multiple Regression for Predicting a Quantitative Target Variable

Shiju Zhang

2026-01-18

Contents

Introduction	1
Load Required Libraries	1
Load and Explore Data	1
Data Preparation	2
Train Multiple Regression Model	2
Make Predictions	3
Calculate Performance Metrics	3
Visualize Results	4
Cross-Validation Results	5
Conclusion	5
Critical Thinking: Why Cross-Validation?	5
Appendix: How Cross-Validation Works	5

Introduction

This document demonstrates a machine learning workflow for multiple regression using a 70/30 train-test split.

Load Required Libraries

```
library(tidyverse)
library(caret)
library(kableExtra)
```

Load and Explore Data

For this example, we'll use the built-in mtcars dataset. In practice, you would load your own dataset.

```
cat("Dataset dimensions:", dim(mtcars)[1], "rows and", dim(mtcars)[2], "columns\n")
## Dataset dimensions: 32 rows and 11 columns
cat("Target variable: mpg (Miles per Gallon)\n")
## Target variable: mpg (Miles per Gallon)
cat("Predictor variables:", paste(names(mtcars)[-1], collapse = ", "), "\n")
## Predictor variables: cyl, disp, hp, drat, wt, qsec, vs, am, gear, carb
```

Data Preparation

```
# Set seed for reproducibility
set.seed(123)

# Create 70/30 train-test split
train_index <- createDataPartition(mtcars$mpg, p = 0.7, list = FALSE)
train_data <- mtcars[train_index, ]
test_data <- mtcars[-train_index, ]

cat("Training set size:", nrow(train_data), "observations\n")

## Training set size: 24 observations
cat("Test set size:", nrow(test_data), "observations\n")

## Test set size: 8 observations
```

Train Multiple Regression Model

```
# Train linear regression model
model <- lm(mpg ~ ., data = train_data) # The dot represents all other variables

# Model summary
summary(model)

## 
## Call:
## lm(formula = mpg ~ ., data = train_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.8335 -1.2145 -0.0044  1.1104  4.1136
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -14.72655  27.33755 -0.539   0.599    
## cyl          1.33019   1.41806  0.938   0.365    
## disp         0.01218   0.01887  0.645   0.530    
## hp          -0.01756   0.02588 -0.678   0.509    
## drat         2.67080   2.51575  1.062   0.308    
## wt          -3.65531   2.37988 -1.536   0.149    
## qsec         1.02957   0.86524  1.190   0.255    
## vs           -0.37848   3.00566 -0.126   0.902    
## am           1.82482   2.74395  0.665   0.518    
## gear         3.74182   2.99973  1.247   0.234    
## carb        -1.42190   1.42397 -0.999   0.336    
## 
## Residual standard error: 2.641 on 13 degrees of freedom
## Multiple R-squared:  0.8971, Adjusted R-squared:  0.818 
## F-statistic: 11.33 on 10 and 13 DF,  p-value: 6.948e-05
```

Make Predictions

```
# Predict on test data
predictions <- predict(model, newdata = test_data)

# Create results data frame
results <- data.frame(
  Actual = test_data$mpg,
  Predicted = predictions
)
head(results)

##          Actual Predicted
## Mazda RX4 Wag     21.0 21.806487
## Duster 360       14.3 13.367084
## Chrysler Imperial 14.7  9.796747
## Dodge Challenger 15.5 17.409006
## Porsche 914-2    26.0 29.354589
## Lotus Europa      30.4 29.017943
```

Calculate Performance Metrics

```
# Calculate metrics
rmse <- sqrt(mean((results$Actual - results$Predicted)^2))
mae <- mean(abs(results$Actual - results$Predicted))
r2 <- cor(results$Actual, results$Predicted)^2
mape <- mean(abs((results$Actual - results$Predicted)/results$Actual)) * 100

# Create metrics table
metrics_table <- data.frame(
  Metric = c("RMSE", "MAE", "R-squared", "MAPE (%)" ),
  Value = c(
    round(rmse, 3),
    round(mae, 3),
    round(r2, 3),
    round(mape, 3)
  ),
  Description = c(
    "Root Mean Square Error",
    "Mean Absolute Error",
    "Coefficient of Determination",
    "Mean Absolute Percentage Error"
  )
)

# Display metrics table
kable(metrics_table, caption = "Model Performance Metrics on Test Set") %>%
  kable_styling(bootstrap_options = c("striped", "hover", "condensed"))
```

Table 1: Model Performance Metrics on Test Set

Metric	Value	Description
RMSE	3.989	Root Mean Square Error
MAE	2.844	Mean Absolute Error

R-squared	0.658	Coefficient of Determination
MAPE (%)	16.631	Mean Absolute Percentage Error


Visualize Results

```
# Create visualization
library(patchwork)

p1 <- ggplot(results, aes(x = Actual, y = Predicted)) +
  geom_point(color = "steelblue", size = 3, alpha = 0.7) +
  geom_abline(intercept = 0, slope = 1, color = "red", linetype = "dashed") +
  labs(title = "Actual vs Predicted Values",
       x = "Actual MPG",
       y = "Predicted MPG") +
  theme_minimal()

p2 <- ggplot(results, aes(y = Predicted - Actual)) +
  geom_boxplot() +
  labs(title = "Distribution of Residuals",
       x = "Residuals (Predicted - Actual)",
       y = "Count") +
  theme_minimal()

p1 + p2
```



Cross-Validation Results

```
# Perform k-fold cross-validation on training data
set.seed(123)
cv_model <- train(mpg ~ .,
                   data = train_data,
                   method = "lm",
                   trControl = trainControl(method = "cv", number = 5))

cv_results <- cv_model$results
kable(cv_results, caption = "5-Fold Cross-Validation Results") %>%
  kable_styling(bootstrap_options = c("striped", "hover", "condensed"))
```

Table 2: 5-Fold Cross-Validation Results

intercept	RMSE	Rsquared	MAE	RMSD	RsquaredSD	MAESD
TRUE	4.390669	0.6540267	3.528197	1.852619	0.4033106	1.64065

Conclusion

The multiple regression model achieved an R-squared of 0.658 on the test set, indicating that approximately 65.8% of the variance in MPG is explained by the model. The RMSE of 3.989 suggests the average prediction error.

Critical Thinking: Why Cross-Validation?

The Problem with Simple Train/Test Split: if you use a 70/30 split (8 vehicles for training, 24 for testing), you might get unlucky.

- What if your test set has mostly vehicles with higher horse power?
- What if your test set has mostly heavy vehicles?
- The model performance could vary wildly depending on which vehicles end up in the test set!

The Solution: Cross-Validation

Cross-validation helps solve this by systematically testing your model on different parts of your data. We creates k folds with original data. We run a model using each of the folds as a test set and the remaining folds as a training set. The performance of each model is averaged to get the final performance metric. The standard deviation of the k separate performance metrics can also be calculated. If we use the RMSE as a performance metric, the final CV performance metric can be expressed as $AVG.RMSE \pm SD$ (say 4.39 ± 0.35), which tells you that

- The typical error is about 4.39 miles per gallon.
- The error varies between 4.04 and 4.74 depending on the data.
- Your model is fairly consistent across different types of vehicles.

Appendix: How Cross-Validation Works

CV takes 3 steps.

- Step 1: Partition the Data We divide the original dataset into k equally sized subsets, called “folds.” For instance, with 100 data points and $k = 5$, we would create 5 folds, each containing 20 data points.

- Step 2: Iterative Training and Testing The algorithm performs k separate modeling iterations. In each iteration:

Test Set: One fold serves as the test set

Training Set: The remaining $(k - 1)$ folds combine to form the training set

This process rotates through all folds, ensuring each data point appears in the test set exactly once.

- Step 3: Performance Aggregation After completing all k iterations, we calculate performance metrics (such as RMSE or accuracy) for each model. The final performance estimate is obtained by averaging these k individual metrics, providing a more robust and reliable assessment of the model's generalization capability.

This methodology offers several advantages:

- (Reduced Variance) By averaging across multiple test sets, we minimize the impact of any particularly favorable or unfavorable data split.
- (Maximized Data Utilization) Every data point contributes to both training and testing phases.
- (Reliable Performance Estimation) The averaged metrics offer a more trustworthy indicator of how the model will perform on unseen data.